

Advanced Message-Passing Programming Exercises

MPI Performance

David Henty

1 Introduction

The purpose of this exercise is to investigate the performance of basic point-to-point MPI operations. You are given a simple ping-pong code that exchanges messages of increasing size between two MPI ranks, and prints out the average time taken and the bandwidth over a large number of repetitions. With this code you can investigate how the following factors affect the performance:

- different underlying protocols for different message sizes;
- different communication hardware (e.g. over the inter-node network or intra-node memory copies);
- different MPI send modes (e.g. synchronous vs. buffered);
- different MPI datatypes (e.g. contiguous or strided data);
- different NUMA regions within a node.

2 Compiling and Running

Check out the IHP.CSS23 repo – `git clone https://github.com/davidhenty/ihpcss2023` – and you should see a `pingpong` directory. The code is contained in `src`.

You should be able to compile it using the Makefile and submit using the supplied Slurm script unchanged. As distributed, the Makefile is set up for Bridges2.

On Bridges2, you will need to load a non-default module:

```
[user@bridges2 ~]$ module load openmpi
```

Compile the code with `make` and submit using `sbatch bridges2.job`.

By default, the code runs on two processes each placed on the same node (so communications will **not** take place over the network), and benchmarks both standard and synchronous sends.

Note that **the program also prints out the exact location of the MPI processes** – which cores of which nodes they are running on – which will enable you to check, for example, whether the two processes are on the same or different nodes. This is important as the messages will only go over the network if you run on different nodes (although the performance within a node is interesting as it should be much faster).

For each mode, two “.plot” files are written containing the times and the bandwidths as a function of message size. The results for the two different modes can be compared using:

```
gnuplot -persist plot_time.gp
gnuplot -persist plot_bandwidth.gp
```

Check that you understand the general form of the graphs before proceeding.

3 Experiments

The supplied gnuplot “.gp” files compare the results for standard and synchronous modes. If you want to compare different modes, or more than two modes, you will have to edit the gnuplot files – the format should be self-explanatory. Also note that the code overwrites the “.plot” files so you will need to make copies after each run if you want to keep a record of the results (e.g. when changing the placement of processes on the cores or nodes).

If you run on more than two processes, the program sends messages between the first (rank zero) and last (rank $N - 1$) processes with all the other processes remaining idle. This is useful when altering the assignment of processes to cores.

By varying `nodes` and `tasks-per-node` in the Slurm script you should have complete control of the placement of the first and last processes.

After having run the default setup – two processes on cores 0 and 1 – you should experiment with the following:

1. Send between different nodes, e.g. run with a total of two processes with one process on each of two nodes.
2. Communication between two processes widely separated on the same node (e.g. `nodes=1` and `tasks-per-node=128`).
3. Can you explain any variations in the latency and bandwidth as you change the process locations?
4. Investigate the performance of buffered or strided sends.